

Selection Pushing in Relational Algebra

In relational algebra, the following equivalencies apply (among others):

$$\sigma_{p_1 \wedge p_2 \wedge \dots \wedge p_n}(R) = \sigma_{p_1}(\sigma_{p_2}(\dots(\sigma_{p_n}(R))\dots)) \quad (1)$$

$$\sigma_p(\sigma_q(R)) = \sigma_q(\sigma_p(R)) \quad (2)$$

$$\sigma_p(R_1 \times R_2) = \sigma_p(R_1) \times R_2 \quad (3)$$

$$\sigma_{R_1.A_1=R_2.A_2}(R_1 \times R_2) = R_1 \bowtie_{R_1.A_1=R_2.A_2} R_2 \quad (4)$$

Some remarks:

- Rule (1) states that a conjunction in a selection predicate may be broken into several nested selections. At the same time, nested selections may be merged into a single selection with a conjunctive predicate.
- Rule (2) states that **nested selections may swap places**.
- Rule (3) states that a **selection can be pushed down** over a cross product, if it only requires the attributes of one of the operands. In the rule as stated above, we assume that predicate p only requires attributes from R_1 . (We need to consult the data dictionary recording the name and attributes of each relation).
- Rule (4) describes how a **selection and a cross product may be merged into a theta join**, provided that the selection predicate is a join condition. This is the case if it compares attributes of R_1 and R_2 .

Implement rule-based selection pushing and perform it on your relational algebra queries. Proceed in these phases:

1. Complex selection predicates are broken up, according to rule (1).
2. All selections are pushed down as far as possible, according to rules (2) and (3).
3. Nested selections are merged again, according to rule (1).
4. Joins are introduced, where possible, according to (4).

Note that there are more rules for the logical optimization of relational algebra, such as rules for join reordering or projection pushing. For *miniHive*, we will make do with this small set of optimization rules (for now).

Write a Python module `raopt.py` that takes a relational algebra query. You may assume the query is the result of the canonical translation of SQL into relational algebra, so it uses only the operators σ , π , ρ and \times).

This is how it should work. The data dictionary `dd` contains the relational schema and can be consulted during selection pushing.

```
>>> import radb.parse
>>> import raopt
>>>
>>> # The data dictionary describes the relational schema.
>>> dd = {}
>>> dd["Person"] = {"name": "string", "age": "integer", "gender": "string"}
>>> dd["Eats"] = {"name": "string", "pizza": "string"}
>>>
>>> stmt = """\project_{Person.name, Eats.pizza}
...         \select_{Person.name = Eats.name}(Person \cross Eats);"""
>>> ra = radb.parse.one_statement_from_string(stmt)
>>>
>>> ra1 = raopt.rule_break_up_selections(ra)
>>> ra2 = raopt.rule_push_down_selections(ra1, dd)
>>> ra3 = raopt.rule_merge_selections(ra2)
>>> ra4 = raopt.rule_introduce_joins(ra3)
>>>
>>> print(ra4)
\project_{Person.name, Eats.pizza} (Person \join_{Person.name = Eats.name} Eats)
```

Remarks: For this project, you are not asked to find any particular optimizations to make your implementation more efficient. Concentrate on a *correct* implementation.

Upload your solution as single file named `raopt.py` to Praktomat (https://praktomat.sdfs.fim.uni-passau.de/SDS_20_WS/), in time for the deadline. `raopt.py` should contain the following methods:

- ```
def rule_break_up_selections(ra): ...
```
- ```
def rule_push_down_selections(ra, dd): ...
```
- ```
def rule_merge_selections(ra): ...
```
- ```
def rule_introduce_joins(ra): ...
```

Praktomat will use the unit tests of `test_raopt.py` and `test_raopt_extended.py` to check your solution. Your solution should also work for other, similar queries as in `test_raopt.py` and `test_raopt_extended.py`. You may start with `test_raopt.py`, and once all tests work, you move on to `test_raopt_extended.py`.

There will be a plagiarism check. It worked really well in last year's course.
